

HW2_P3P4

October 9, 2022

```
[1]: """  
      Edward Pascual-Bautista  
      ECGR 4105 HW2  
      """  
import numpy as np  
import matplotlib.pyplot as plt  
import pandas as pd
```

```
[2]: from sklearn.datasets import load_breast_cancer  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler
```

```
[3]: breast = load_breast_cancer()  
  
breast_data = breast.data  
breast_data.shape
```

[3]: (569, 30)

```
[4]: breast_input = pd.DataFrame(breast_data)  
breast_input.head()
```

```
[4]:
```

	0	1	2	3	4	5	6	7	8	\		
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419			
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812			
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069			
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597			
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809			
		9	...	20	21	22	23	24	25	26	27	\
0	0.07871	...	25.38	17.33	184.60	2019.0	0.1622	0.6656	0.7119	0.2654		
1	0.05667	...	24.99	23.41	158.80	1956.0	0.1238	0.1866	0.2416	0.1860		
2	0.05999	...	23.57	25.53	152.50	1709.0	0.1444	0.4245	0.4504	0.2430		
3	0.09744	...	14.91	26.50	98.87	567.7	0.2098	0.8663	0.6869	0.2575		
4	0.05883	...	22.54	16.67	152.20	1575.0	0.1374	0.2050	0.4000	0.1625		
		28		29								
0	0.4601		0.11890									

```
1  0.2750  0.08902
2  0.3613  0.08758
3  0.6638  0.17300
4  0.2364  0.07678
```

[5 rows x 30 columns]

```
[5]: breast_labels = breast.target
breast_labels.shape
```

[5]: (569,)

```
[6]: labels = np.reshape(breast_labels,(569,1))
final_breast_data = np.concatenate([breast_data,labels],axis = 1)
final_breast_data.shape
```

[6]: (569, 31)

```
[7]: breast_dataset = pd.DataFrame(final_breast_data)
breast_dataset
```

```
[7]:
```

	0	1	2	3	4	5	6	7	8	\
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	0.2419	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	0.1812	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	0.2069	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	0.2597	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	0.1809	
..	
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	0.1726	
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	0.1752	
566	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	0.1590	
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	0.2397	
568	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	0.1587	
	9	...	21	22	23	24	25	26	27	\
0	0.07871	...	17.33	184.60	2019.0	0.16220	0.66560	0.7119	0.2654	
1	0.05667	...	23.41	158.80	1956.0	0.12380	0.18660	0.2416	0.1860	
2	0.05999	...	25.53	152.50	1709.0	0.14440	0.42450	0.4504	0.2430	
3	0.09744	...	26.50	98.87	567.7	0.20980	0.86630	0.6869	0.2575	
4	0.05883	...	16.67	152.20	1575.0	0.13740	0.20500	0.4000	0.1625	
..	
564	0.05623	...	26.40	166.10	2027.0	0.14100	0.21130	0.4107	0.2216	
565	0.05533	...	38.25	155.00	1731.0	0.11660	0.19220	0.3215	0.1628	
566	0.05648	...	34.12	126.70	1124.0	0.11390	0.30940	0.3403	0.1418	
567	0.07016	...	39.42	184.60	1821.0	0.16500	0.86810	0.9387	0.2650	
568	0.05884	...	30.37	59.16	268.6	0.08996	0.06444	0.0000	0.0000	

	28	29	30
0	0.4601	0.11890	0.0
1	0.2750	0.08902	0.0
2	0.3613	0.08758	0.0
3	0.6638	0.17300	0.0
4	0.2364	0.07678	0.0
..
564	0.2060	0.07115	0.0
565	0.2572	0.06637	0.0
566	0.2218	0.07820	0.0
567	0.4087	0.12400	0.0
568	0.2871	0.07039	1.0

[569 rows x 31 columns]

```
[8]: features = breast.feature_names

feature_labels = np.append(features, 'label')

breast_dataset.columns = feature_labels

breast_dataset.head()
```

```
[8]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
0	0.07871	...	17.33	184.60	2019.0	
1	0.05667	...	23.41	158.80	1956.0	
2	0.05999	...	25.53	152.50	1709.0	
3	0.09744	...	26.50	98.87	567.7	
4	0.05883	...	16.67	152.20	1575.0	

	worst smoothness	worst compactness	worst concavity	worst concave points	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	

2	0.1444	0.4245	0.4504	0.2430
3	0.2098	0.8663	0.6869	0.2575
4	0.1374	0.2050	0.4000	0.1625

	worst symmetry	worst fractal dimension	label
0	0.4601	0.11890	0.0
1	0.2750	0.08902	0.0
2	0.3613	0.08758	0.0
3	0.6638	0.17300	0.0
4	0.2364	0.07678	0.0

[5 rows x 31 columns]

```
[9]: """breast_dataset['label'].replace(0, 'Benign',inplace=True)
breast_dataset['label'].replace(1, 'Malignant',inplace=True)

breast_dataset.head()"""
```

```
[9]: "breast_dataset['label'].replace(0,
'Benign',inplace=True)\nbreast_dataset['label'].replace(1,
'Malignant',inplace=True)\n\nbreast_dataset.head()"
```

```
[10]: X = breast.data
y = breast.target
```

```
[11]: X.shape
```

```
[11]: (569, 30)
```

```
[12]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.8,
↳test_size = 0.2, random_state = 0)
X_train.shape
```

```
[12]: (455, 30)
```

```
[13]: sc = StandardScaler()

X_train_std = sc.fit_transform(X_train)
X_test_std = sc.transform(X_test)
```

```
[14]: from sklearn.linear_model import LogisticRegression

classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train_std, y_train)
```

```
[14]: LogisticRegression(random_state=0)
```

```
[15]: y_pred = classifier.predict(X_test_std)
```

```
[16]: from sklearn import metrics
```

```
print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))
print("Precision: ", metrics.precision_score(y_test, y_pred))
print("Recall: ", metrics.recall_score(y_test, y_pred))
```

```
Accuracy: 0.9649122807017544
Precision: 0.9701492537313433
Recall: 0.9701492537313433
```

```
[17]: from sklearn.metrics import confusion_matrix
```

```
cnf_matrix = confusion_matrix(y_test, y_pred)
cnf_matrix
```

```
[17]: array([[45,  2],
           [ 2, 65]], dtype=int64)
```

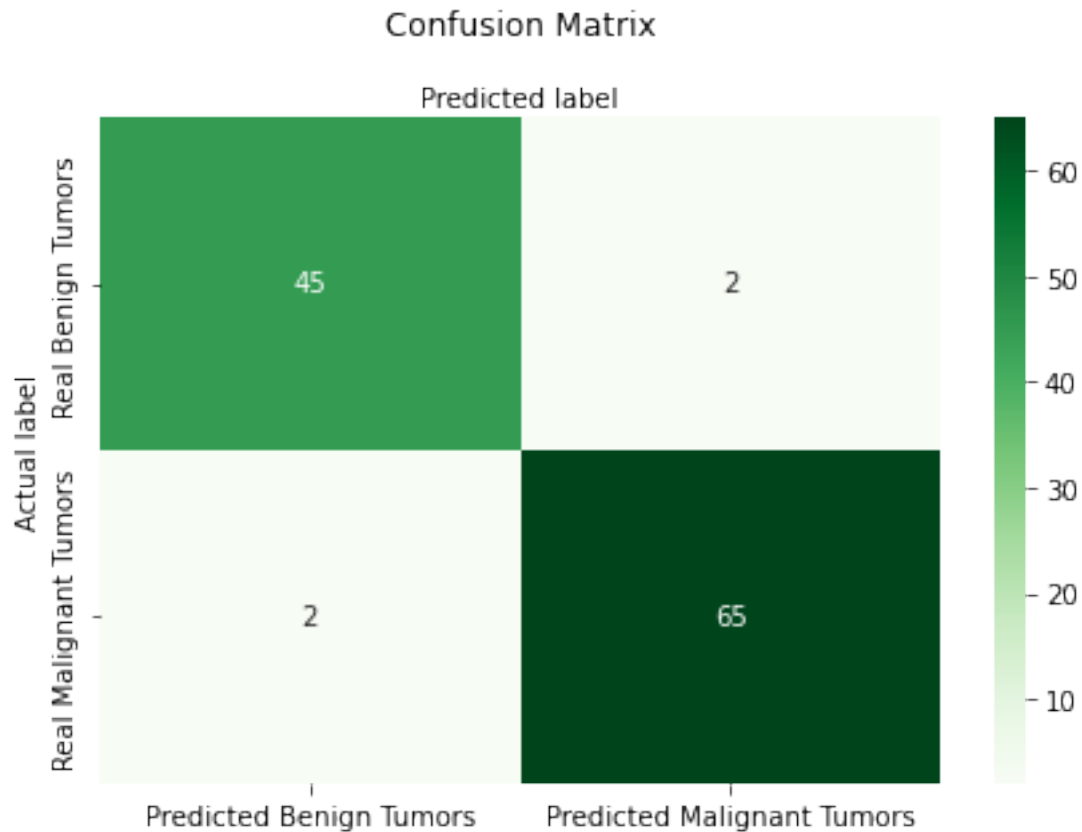
```
[18]: import seaborn as sns
```

```
class_names = [0,1]
```

```
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
```

```
sns.heatmap(pd.DataFrame(cnf_matrix), annot = True, cmap = plt.cm.Greens, fmt = 'g',
            yticklabels = ["Real Benign Tumors", "Real Malignant Tumors"],
            xticklabels = ["Predicted Benign Tumors", "Predicted Malignant Tumors"])
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion Matrix', y = 1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
[18]: Text(0.5, 257.44, 'Predicted label')
```



```
[19]: classifier_weight = LogisticRegression(penalty='l1', class_weight = '
      ↪'balanced', solver = 'liblinear', random_state = 0)
      classifier_weight.fit(X_train_std, y_train)
```

```
[19]: LogisticRegression(class_weight='balanced', penalty='l1', random_state=0,
      solver='liblinear')
```

```
[20]: y_pred = classifier_weight.predict(X_test_std)
```

```
[21]: print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))
      print("Precision: ", metrics.precision_score(y_test,y_pred))
      print("Recall: ", metrics.recall_score(y_test, y_pred))
```

Accuracy: 0.956140350877193

Precision: 0.9558823529411765

Recall: 0.9701492537313433

```
[22]: from sklearn.model_selection import KFold
      from sklearn.model_selection import cross_val_score
```

```
[23]: kfold = KFold(n_splits = 5, random_state = 0, shuffle = True)

model = LogisticRegression(solver = 'liblinear')

results = cross_val_score(model, X, y, cv = kfold)

print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 95.434% (2.737%)

```
[24]: kfold = KFold(n_splits = 10, random_state = 0, shuffle = True)

model = LogisticRegression(solver = 'liblinear')

results = cross_val_score(model, X, y, cv = kfold)

print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 95.432% (3.858%)

```
[25]: kfold = KFold(n_splits = 5, random_state = 0, shuffle = True)

model = LogisticRegression(penalty = 'l1', class_weight = 'balanced', solver = 'liblinear')

results = cross_val_score(model, X, y, cv = kfold)

print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 95.435% (2.680%)

C:\Users\MexDrakus\anaconda3\lib\site-packages\sklearn\svm_base.py:1206:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
warnings.warn(

```
[26]: kfold = KFold(n_splits = 10, random_state = 0, shuffle = True)

model = LogisticRegression(penalty = 'l1', class_weight = 'balanced', solver = 'liblinear')

results = cross_val_score(model, X, y, cv = kfold)

print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

Accuracy: 95.432% (3.695%)

```
[ ]:
```